


Stream Processing with Kafka and Python

What is it about?

- Context: data engineering
- Goal: make decision process less intuitive and more informed
- Approach: datawarehouse
 - Collect information in one place, process on schedule
 - Make well-informed decisions using all of it
- Approach: streaming
 - Processing as a reaction to an event
 - Make decisions as soon as possible

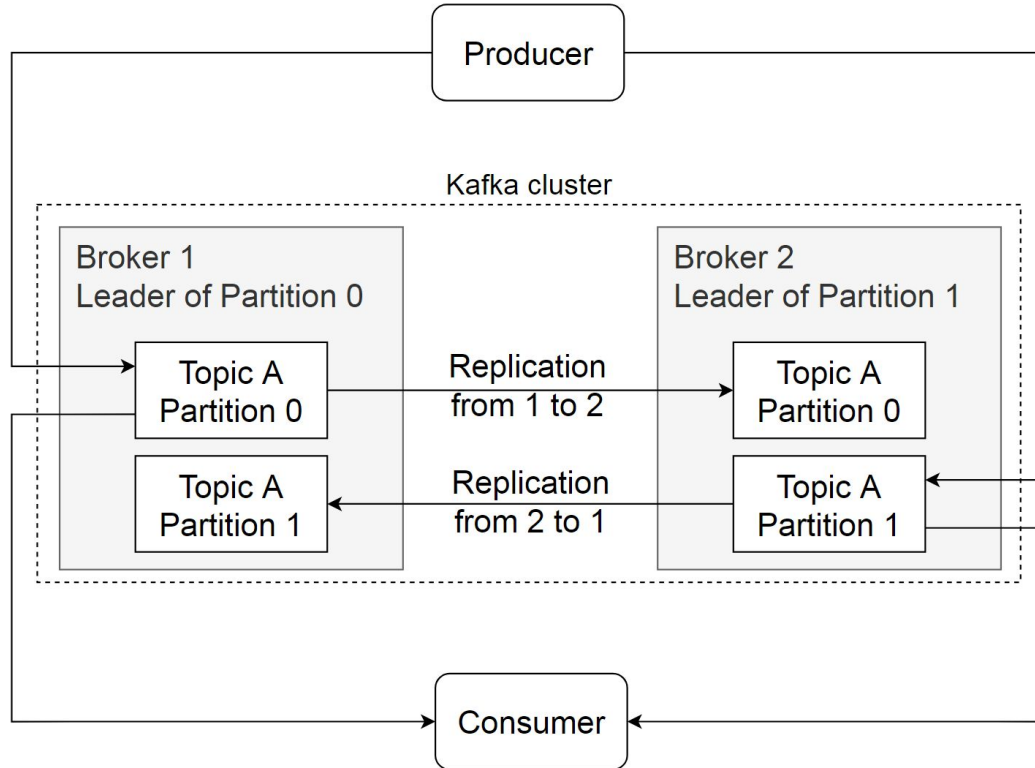
The story of data



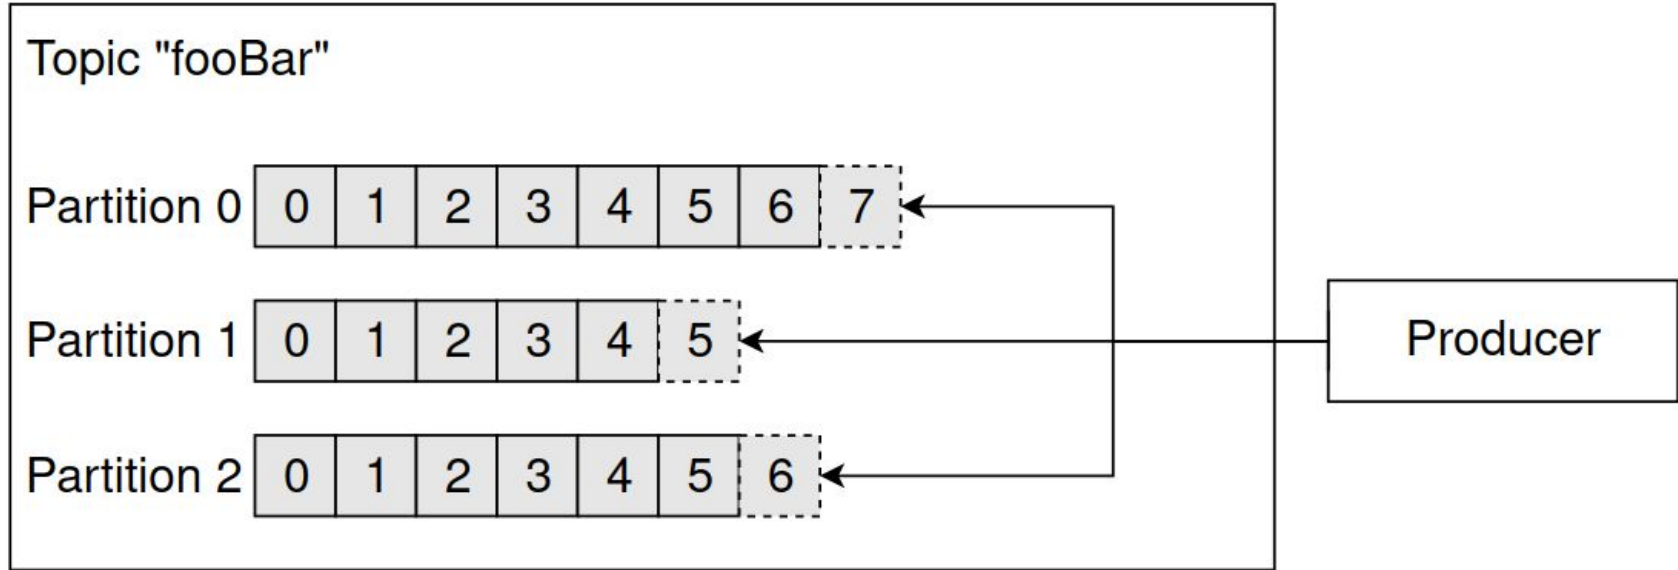
Welcome to Kafka

- Publish/subscribe messaging system
- Unit of data is a *message*
 - Similar to a row in relational databases
- Message belongs to a *topic*
 - Similar to a table in relational databases
- Topic is split in *partitions*


Why Kafka is called distributed?




Writing to Kafka



Reading from Kafka



How to write messages to files *and keep your mental health in shape*



Configuration: topic

- `num.partitions`
 - Decide basing on future throughput vs consumer capacity
- `log.retention.ms` and `log.retention.bytes`
- `log.segment.bytes` and `log.segment.ms`
 - Retention happens per segments
- `message.max.bytes`
 - Defaults to 1 Mb; beware of performance impact
 - Adjust `fetch.message.max.bytes` and `replica.fetch.max.bytes`

Configuration: producer

- `acks` configures desired delivery guarantee
 - `0` means fire and forget
 - `1` means leader replica received the message
 - `all` means all in-sync replicas received the message
- `batch.size` configures a tradeoff between latency and throughput
 - Large batches increase maximum of messages per second
 - Small batches decrease latency for a single message
- `linger.ms` is the amount of time to wait before sending the batch
 - Defaults to 0; use higher value to increase throughput

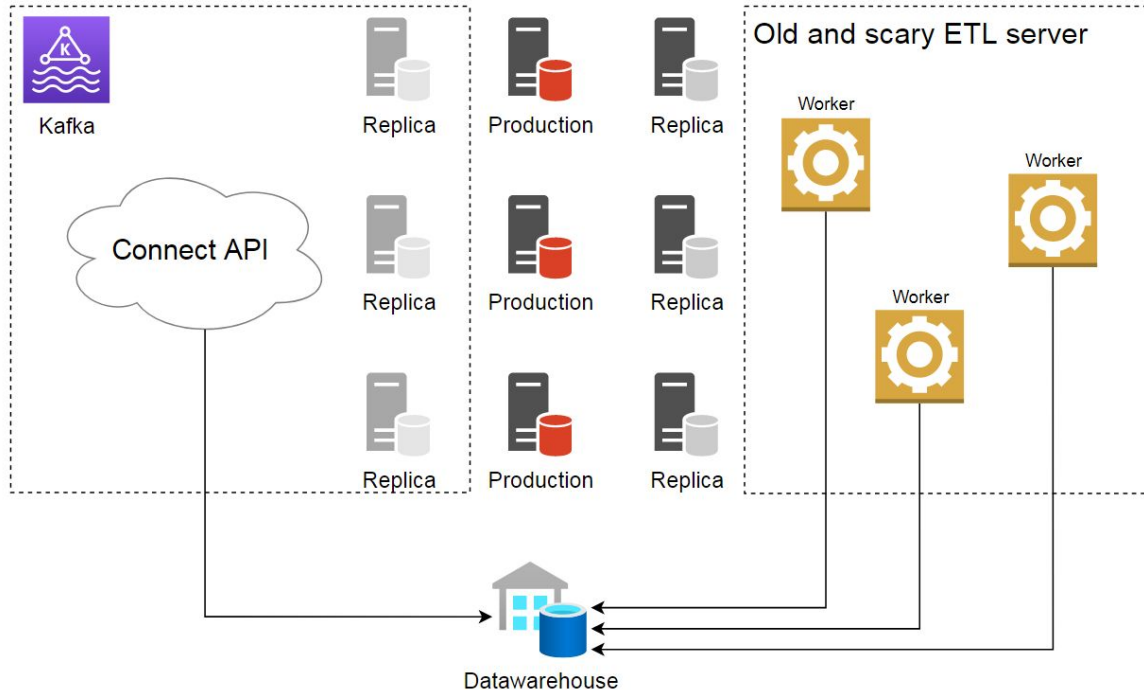
Configuration: consumer

- `auto.offset.reset` controls handling of a new partition
 - `latest` is the default option to read from newest records
 - `earliest` means reading from the oldest records
- `enable.auto.commit` can be `true` or `false`
- `partition.assignment.strategy` defines how partitions are assigned to consumers
 - `Range` strategy is the default one
 - `RoundRobin` strategy provides more evenly distributed assignments
 - Custom implementations are possible

Extending Kafka: Kafka Connect

- Producer/Consumer APIs require programming
 - Connecting Kafka with typical data sources should be easier
- Kafka Connect is a tool for connecting Kafka with other systems
 - Supports two types of connectors: **source** and **sink**
 - Runs in distributed or standalone mode
 - Uses Kafka itself as a backend to keep state
 - Provides REST API to create connectors

Creating of near real time datawarehouse *or how to trick your DBA with a fake replication server*



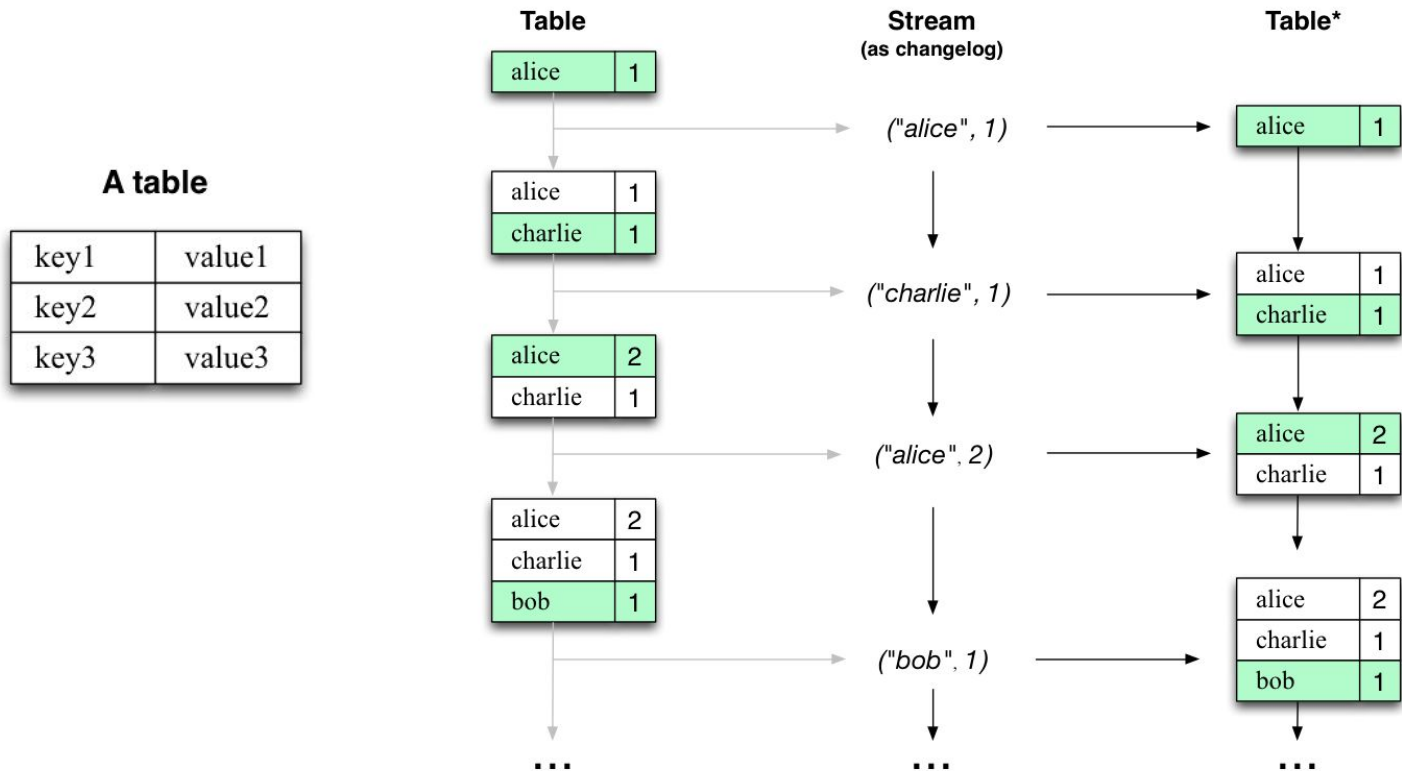
Extending Kafka: Kafka Streams

- Once you want to do data processing, you need to consume it first
 - Do you?
 - What if all data you need is *already in Kafka*?
- Kafka Streams allows to run data processing inside Kafka
- Write Streams app in Java and let Kafka handle operational load


Kafka Streams: Define time

- Event time
 - Created “at the source”
- Log append or ingestion time
 - When message is written into a topic partition
- Processing time
 - When message is being consumed by Streams app
 - It can be milliseconds or hours and days greater than event time
 - Not really reliable

Kafka Streams: Stream-Table duality



Kafka Streams: Processor Topology



PROCESSOR TOPOLOGY

- Read from Source processor
- Write into Sink processor
- Transform in Stream processor
 - Streams DSL is available in Java: `map`, `filter`, `join`
 - Aggregations are also possible
 - Windowing with a grace period
 - State stores: persistent key-value or in-memory hashmap or else

Aggregating Wikipedia events: demo app

<https://github.com/ilia-khaustov/sytac-devjam2021-wikiapp>

